

# Class 07: Machine Learning

Lana (PID:A17013518)


## Clustering Methods

The broad goal here is to find grouping (clusters) in your input data.

### Kmeans

First, let's make up some data to cluster.

```
x <- rnorm(1000)  
hist(x)
```



Make a vector of length 60 with 30 points centered at -3 and 30 points centered at +3

```
tmp <- c(rnorm(30, mean = -3), rnorm(30, mean=3))
tmp
```


```
[1] -4.2636379 -4.2655327 -3.2921880 -3.1532843 -2.0853496 -3.3542151
[7] -2.6278981 -3.5762702 -3.8662051 -3.5534463 -2.7073550 -4.2948711
[13] -2.5434864 -3.8787771 -2.8245481 -3.1000712 -2.3996220 -5.4969447
[19] -3.6359823 -3.4217302 -2.5206616 -3.6814908 -3.3113487 -3.6836124
[25] -0.9114059 -3.5052324 -5.3590019 -2.8990608 -0.9126699 -2.8240840
[31]  2.9669712  3.2297234  1.7828317  2.9304934  2.7824536  3.2346053
[37]  2.9082036  2.4601156  2.0678986  2.6143794  2.8871483  4.5250711
[43]  3.8571460  2.8892992  4.2444765  4.0408944  2.4964628  3.2191740
[49]  3.8705715  4.3421666  2.5293236  2.0007628  5.1667340  0.6551775
[55]  2.9369875  2.4235691  3.4345135  4.6592965  1.4338930  4.8880406
```

I will now make a wee x and y dataset with 2 groups of points.

```
rev(c(1:5))
```

```
[1] 5 4 3 2 1
```

```
x <- cbind(x=tmp, y=rev(tmp))
plot(x)
```



```
k <- kmeans(x, centers = 2)
k
```

K-means clustering with 2 clusters of sizes 30, 30

Cluster means:

	x	y
1	3.115946	-3.264999
2	-3.264999	3.115946

Clustering vector:

```
[1] 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 1 1 1 1 1 1 1 1
[39] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1
```

Within cluster sum of squares by cluster:

```
[1] 61.33383 61.33383
(between_SS / total_SS = 90.9 %)
```

Available components:

[1] "cluster"	"centers"	"totss"	"withinss"	"tot.withinss"
[6] "betweenss"	"size"	"iter"	"ifault"	

Q. From your result object `k` how many points are in each cluster?

```
k$size
```

```
[1] 30 30
```

Q. What “component” of your result object details the cluster membership?

```
k$cluster
```

```
[1] 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 1 1 1 1 1 1 1 1
[39] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1
```


Q. Cluster centers?

```
k$centers
```

	x	y
1	3.115946	-3.264999
2	-3.264999	3.115946


Plot for clustering results

```
plot(x, col = k$cluster)
points(k$centers, col = "pink", pch = 15, cex = 2)
```



We can cluster into 4 groups

```
# kmeans  
k4 <- kmeans(x, centers = 4)  
# plot results  
plot(x, col = k4$cluster)
```



A big limitation of `kmeans` is that it does what you ask even if you ask for silly clusters.

## Hierarchical Clustering

The main base R function for Hierarchical Clustering is `hclust()`. Unlike `kmeans()` you can not just pass it your data as input. You first need to calculate distance matrix.

```
d <- dist(x)
hc <- hclust(d)
hc
```


Call:

```
hclust(d = d)
```

```
Cluster method : complete
Distance       : euclidean
Number of objects: 60
```

Use `plot()` to view results

```
plot(hc)
abline(h=10, col = "pink")
```




To make the “cut” and get our cluster membership vector we can use the `cutree()` function.

```
grps <- cutree(hc, h=10)
grps
```

```
[1] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 2 2 2 2 2 2 2 2
[39] 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2
```

Make a plot of data colored by hclust results

```
plot(x, col=grps)
```



## principal Component Analysis (PCA)

Here we will do principal component analysis (PCA) on some food data from the UK.

```
url <- "https://tinyurl.com/UK-foods"
x <- read.csv(url, row.names = 1)
```

```
# rownames(x) <- x[,1]
# x <- x[, -1]
# x
```

Q1. How many rows and columns are in your new data frame named x? What R functions could you use to answer this questions?

```
dim(x)
```

```
[1] 17  4
```


Q2. Which approach to solving the ‘row-names problem’ mentioned above do you prefer and why? Is one approach more robust than another under certain circumstances?



```
# I prefer the first approach just because it looks more simple and clean in the code and
```


Q3: Changing what optional argument in the above `barplot()` function results in the following plot?

```
# Changing beside to false results in changing the barplot because if false, columns of height
barplot(as.matrix(x), beside=F, col=rainbow(nrow(x)))
```



Q5: Generating all pairwise plots may help somewhat. Can you make sense of the following code and resulting figure? What does it mean if a given point lies on the diagonal for a given plot?

```
pairs(x, col=rainbow(10), pch=16)
```



Q6. What is the main differences between N. Ireland and the other countries of the UK in terms of this data-set?

## PCA to the Rescue

The main “base” R function of PCA is called `prcomp()`.


```
pca <- prcomp(t(x))
summary(pca)
```

Importance of components:

	PC1	PC2	PC3	PC4
Standard deviation	324.1502	212.7478	73.87622	2.921e-14
Proportion of Variance	0.6744	0.2905	0.03503	0.000e+00
Cumulative Proportion	0.6744	0.9650	1.00000	1.000e+00


Q7. Complete the code below to generate a plot of PC1 vs PC2. The second line adds text labels over the data points.

```
plot(pca$x[,1], pca$x[,2], xlab="PC1", ylab="PC2", xlim=c(-270,500))
text(pca$x[,1], pca$x[,2], colnames(x))
```



Q8. Customize your plot so that the colors of the country names match the colors in our UK and Ireland map and table at start of this document.

```
plot(pca$x[,1], pca$x[,2], xlab="PC1", ylab="PC2", xlim=c(-270,500))
text(pca$x[,1], pca$x[,2], colnames(x))
```



Q. How much variance is captured in 2 PCs

96.5%

To make our main “PC score plot” or “PC1 vs PC2 plot”, or “PC plot” or “ordination plot”

```
attributes(pca)
```

```
$names
```

```
[1] "sdev"      "rotation" "center"    "scale"     "x"
```

```
$class
```


```
[1] "prcomp"
```

We are after the `pca$x` result component to make our main PCA plot.

```
pca$x
```

	PC1	PC2	PC3	PC4
England	-144.99315	-2.532999	105.768945	-9.152022e-15
Wales	-240.52915	-224.646925	-56.475555	5.560040e-13
Scotland	-91.86934	286.081786	-44.415495	-6.638419e-13
N.Ireland	477.39164	-58.901862	-4.877895	1.329771e-13

```
mycols <- c("orange", "red", "blue", "darkgreen")
plot(pca$x[,1], pca$x[,2], col=mycols, pch=16, xlab = "PC1 (67.4%)", ylab = "PC2 (29%)")
```



Another important result from PCA is how the original variables (in this case foods) contribute to the PCs.

This is contained in the `pca$rotation` object - folks often call this the “loading” or “contributions” to the PCs

```
pca$rotation
```

	PC1	PC2	PC3	PC4
Cheese	-0.056955380	0.016012850	0.02394295	-0.409382587
Carcass_meat	0.047927628	0.013915823	0.06367111	0.729481922
Other_meat	-0.258916658	-0.015331138	-0.55384854	0.331001134
Fish	-0.084414983	-0.050754947	0.03906481	0.022375878
Fats_and_oils	-0.005193623	-0.095388656	-0.12522257	0.034512161
Sugars	-0.037620983	-0.043021699	-0.03605745	0.024943337
Fresh_potatoes	0.401402060	-0.715017078	-0.20668248	0.021396007
Fresh_Veg	-0.151849942	-0.144900268	0.21382237	0.001606882
Other_Veg	-0.243593729	-0.225450923	-0.05332841	0.031153231

Processed_potatoes	-0.026886233	0.042850761	-0.07364902	-0.017379680
Processed_Veg	-0.036488269	-0.045451802	0.05289191	0.021250980
Fresh_fruit	-0.632640898	-0.177740743	0.40012865	0.227657348
Cereals	-0.047702858	-0.212599678	-0.35884921	0.100043319
Beverages	-0.026187756	-0.030560542	-0.04135860	-0.018382072
Soft_drinks	0.232244140	0.555124311	-0.16942648	0.222319484
Alcoholic_drinks	-0.463968168	0.113536523	-0.49858320	-0.273126013
Confectionery	-0.029650201	0.005949921	-0.05232164	0.001890737

We can make a plot along PC1.

```
library(ggplot2)

contrib <- as.data.frame(pca$rotation)

ggplot(contrib) +
  aes(PC1, rownames(contrib)) +
  geom_col(col = "pink")
```

